

LOUSTRINI: A Lustre Model Checker using (H-)Houdini Invariant Learning Algorithm Or me learning about invariant learning algorithms

Rémi Germe

09.01.2026

Contents

First experiments with SMT solvers	2
Translating Lustre to SMT expressions	4
Invariants learning algorithms, Houdini, H-Houdini	9
Future work & conclusion	18

First experiments with SMT solvers

First experiments with SMT solvers

Goal: SMT-solver-agnostic model checker \Rightarrow use a frontend library \Rightarrow in OCaml: `Smt.ml` [1]

- Surprising behavior from `Smt.ml`, played with different solvers as well as the SMTLIB2 format directly (and was stuck in a dependency hell regarding `opam`, `z3`, `llvm`, `ld`, ...)
-

Aftermath of these initial experiments:

- discovered unsound behavior in Bitwuzla mappings of `Smt.ml` (see [issue #465](#))
 - discovered bug in AltErgo [2] support of `Smt.ml` (see [discussion #450](#))
 - more generally, clarification of the current limitations of `Smt.ml` (also [discussion #450](#))
 - re-discovered a known issue in AltErgo failing to conclude SAT (see [issue #1323](#))
-

\Rightarrow Loustrini is *not* solver-agnostic (using Z3 [3] - best would have been cvc5 [4] for *abducts*).

Translating Lustre to SMT expressions

Encoding à la k-induction

Encoding presented in the handout.

$\Delta(n)$ encodes the system equations at time n .

Checking a property P is inductive on system Δ :

$$\text{(initiation)} \quad \Delta(0) \Rightarrow P(0)$$

$$\text{(consecution)} \quad \Delta(n) \wedge \Delta(n+1) \wedge P(n) \Rightarrow P(n+1)$$

Remarks

Wrong semantics for \rightarrow

I implemented a wrong semantics for $e \rightarrow e'$:

$$1 \rightarrow 2 \rightarrow 3 \text{ generates } \begin{cases} 1 \text{ if } n=0 \\ 2 \text{ if } n=1 \\ 3 \text{ if } n \geq 2 \end{cases} \text{ instead of } \begin{cases} 1 \text{ if } n=0 \\ 3 \text{ if } n \geq 1 \end{cases}.$$

Properties referring to n and $n + 1$

They require careful thinking:

$$\underbrace{\Delta(n) \wedge \Delta(n+1)}_{\text{do not constrain } x(n+2)} \wedge P(n) \Rightarrow \underbrace{P(n+1)}_{\text{refers to } x(n+2)} ? \text{ will fail, so still sound}$$

\Rightarrow would require to strengthen our lhs with $\Delta(n+2)$

I did **not** consider those properties.

Encoding as a transition system

Transition system (I, T) using state variables and primed variables, verifying a property P:

$$I \Rightarrow P \text{ and } P \wedge T \Rightarrow P'$$

Handling pre e : introduce a new state variable $S_{\text{id}}^{\text{pre}} = \{\text{init} = \emptyset; \text{next} = \llbracket e \rrbracket\}$

Handling of $e \rightarrow e'$: use $\text{ite}(S_i^\rightarrow, \llbracket e \rrbracket, \llbracket e' \rrbracket)$ with:

- $S_0^\rightarrow = \{\text{init} = \text{true}; \text{next} = \text{false}\}$
- $S_1^\rightarrow = \{\text{init} = \text{false}; \text{next} = S_0^\rightarrow\}$
- $S_2^\rightarrow = \{\text{init} = \text{false}; \text{next} = S_1^\rightarrow\}$
- ...

We also need to make sure we have **at most one** of the S_i^\rightarrow to be true (+ harder to obtain positive examples).



$$\Delta(n) \wedge \Delta(n+1) \wedge P(n) \Rightarrow P(n+1)$$

Even less precise, the consecution might fail because of **spurious counterexamples** (unreachable situations), already the case before but less likely.

Tuples.

Treat a n-tuple as n-expressions and translate each member separately.

Node calls.

Instantiation (inlining) of nodes at call site.

How to learn efficiently invariants for a node (and not just for its instances)?

In real-world projects (see [discussion #1256](#) and [5], [6]):

- compositional reasoning
- modular reasoning
- progressive refinements
- ...

This issue will **not** be addressed in the project.

Invariants learning algorithms, Houdini, H-Houdini

Overview of invariant learning algorithms

Goal: prove a safety property P . Let's not use *k-induction* but *invariant learning* instead.

Challenge: find a property H such that:

$$\text{(initiation)} \quad I \Rightarrow H$$

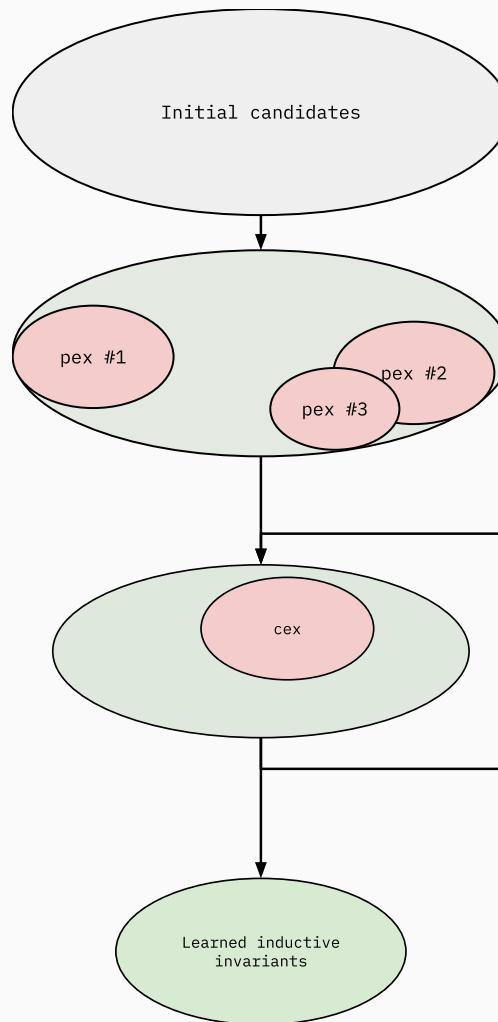
$$\text{(consecution)} \quad H \wedge T \Rightarrow H'$$

$$\text{(implies desired property)} \quad H \Rightarrow P$$

How to find such an H ?

Bottom-up approaches	Top-down approaches (property-directed)
Learn as many invariants as possible	Learn relevant invariants to prove a given desired property
<ul style="list-style-type: none">• Houdini [7]• Instantiation-based invariant discovery [6]	<ul style="list-style-type: none">• IC3 [8]• H-Houdini [9]

Houdini: overview [6], [7]



- Generating a lot of candidates using templates (see next slide)
- Initial sift: removing all candidates that do no satisfy a trace of execution (see next next slide)
- Inductivity sift: removing all candidates that break inductivity, and iterate until reaching an inductive set (see next next slide)
- Final learned inductive invariants

Houdini: generating invariants

Variables below are all evaluated at n (n and $n - 1$ (+ init to true) would have been possible).

Booleans. $\mathcal{I}_b ::= b = \text{true} \mid b = \text{false} \mid b_1 = b_2 \mid b_1 = \neg b_2$

Integers. $\mathcal{I}_i ::= i \diamond \text{cst} \mid i_1 \diamond i_2$

Reals. Same as for integers.

Where:

- $\text{cst} \in \{0, 1, -1\} \cup \{\text{constants of interest (hardcoded in the program)}\}$
- $\diamond ::= \geq \mid > \mid \leq \mid < \mid = \mid \neq$

We obtain a set of candidates $H = \bigwedge_{i=1}^d h_i$.

SOUND but absolutely not **COMPLETE**

but Houdini is complete *relative* to the templates

Houdini: sifting candidates

Inductivity check.

$$\Delta(n) \wedge \Delta(n+1) \wedge H(n) \wedge \neg H(n+1) ?$$

UNSAT : H is inductive OR $H(n)$ contradicts $\Delta(n) \wedge \Delta(n+1)$

SAT : \exists cex, $\begin{cases} \text{cex} \models H(n) \\ \text{cex} \models \neg H(n+1) \end{cases}$ i.e. $\exists i \in \{1, \dots, d\}$, $\underbrace{\text{cex} \models \neg h_i(n+1)}_{\text{remove all these candidates}}$

Better than checking each h_i separately.

Initial sift.

1. Prune the search space by removing many candidates that do not hold.
2. Ensure H is non-contradictory (vacuously false).

$$\Delta(0) \wedge \dots \wedge \Delta(k) \wedge \neg H(k)$$

UNSAT : H is consistent with step k

SAT : \exists cex, ... (similar as above, and we keep iterating for this k)

Houdini: demo

Extra-prunning of “obvious” invariants.

ic3.lus

```
x = 1 -> pre x + 1;  
y = 1 -> pre (x + y);  
ok = y >= 0;
```

Not k-inductive!

fib.lus

With/without explicit state variables.

Limitation: we cannot learn an invariant not present in the templates.

H-Houdini: overview [9]

Motivation.

- candidates generation is subject to combinational explosion,
- sifting may require *many, large* SMT queries.

H-Houdini (“Hierarchical Houdini”) aims at solving these by making Houdini **property-directed**.

```
def h-houdini(p_target):
    V = SLICE(p_target) # extract relevant variables
    P = MINE(p_target, V) # generate candidates
    while # flag to keep iterating:
        A = ABDUCT(p_target, P) # extract a proposition that fixes (makes inductive) p_target
        if A is None: return None
        H = p_target
        for p in A:
            h_sol = h-houdini(p)
            if h_sol is None: # break, keep iterating the while loop
                H = H ^ h_sol
    return H
```

python

Slicing.

- ▷ $\text{SLICE}(p_{\text{target}})$ extracts the variables that influence the inductivity of p_{target}
- ▷ implemented by instrumenting the Lustre → SMT translation

Mining.

- ▷ $\text{MINE}(p_{\text{target}}, V)$ generates invariants according to some templates
- ▷ implemented by reusing functions for Houdini - not ideal as they do not take p_{target} as an input and so subsequent sifting is performed

H-Houdini: abducting

`ABDUCT(p_target, P)` returns `None` or a list of predicates A derived from P that fixes p_{target} :
 $A \wedge p_{\text{target}} \Rightarrow p_{\text{target}}'$. If called multiple times, it returns a different abduct each time.

The paper proposes the following method:

$$\bigwedge_i P_i \wedge p_{\text{target}} \wedge \neg p_{\text{target}}'$$

UNSAT : we extract a minimal unsat core A

SAT : no possible abduct using P

But I fail to see how this method can generate several *different* abducts.

Using SMT solvers:

- ▷ to the best of my knowledge, Z3 does not provide any builtin method
- ▷ cvc5 has a builtin `getAbductNext()` method but no official OCaml bindings

Also: Z3 `get_unsat_core` returns a list, but in practice it is only one element which is a big conjunct (not handy).

The implementation is currently bugged.

Future work & conclusion

Future work & conclusion

Future work

“Minor” improvements for H-Houdini:

- mining according to p_{target}
- topological ordering of equations for slicing
- memoization

Major issue: abducting.

Key insights

- bottom-up vs top-down (property-directed) approaches
- limited to templates (see IC3)
- in real world: used as complement to other techniques
- in real world: modular reasoning

References

References

- [1] J. M. Pereira *et al.*, “Smt.ml: A Multi-Backend Frontend for SMT Solvers in OCaml,” 2024.
- [2] S. Conchon, A. Coquereau, M. Iguernlala, and A. Mebsout, “Alt-Ergo 2.2,” in *SMT Workshop: International Workshop on Satisfiability Modulo Theories*, 2018.
- [3] L. De Moura and N. Bjørner, “Z3: an efficient SMT solver,” in *Proceedings of the Theory and Practice of Software, 14th International Conference on Tools and Algorithms for the Construction and Analysis of Systems*, in TACAS'08/ETAPS'08. Budapest, Hungary: Springer-Verlag, 2008, pp. 337–340.
- [4] H. Barbosa *et al.*, “cvc5: A Versatile and Industrial-Strength SMT Solver,” in *Tools and Algorithms for the Construction and Analysis of Systems - 28th International Conference, TACAS 2022, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2022, Munich, Germany, April 2-7, 2022, Proceedings, Part I*, D. Fisman and G. Rosu, Eds., in Lecture Notes in Computer Science, vol. 13243. Springer, 2022, pp. 415–442. doi: 10.1007/978-3-030-99524-9_24.
- [5] A. Champion, A. Mebsout, C. Sticksel, and C. Tinelli, “The Kind 2 model checker,” in *International Conference on Computer Aided Verification*, 2016, pp. 510–517.
- [6] T. Kahrilas, Y. Ge, and C. Tinelli, “Instantiation-Based Invariant Discovery,” in *NASA Formal Methods*, M. Bobaru, K. Havelund, G. J. Holzmann, and R. Joshi, Eds., Berlin, Heidelberg: Springer Berlin Heidelberg, 2011, pp. 192–206.
- [7] C. Flanagan and K. R. M. Leino, “Houdini, an Annotation Assistant for ESC/Java,” in *Proceedings of the International Symposium of Formal Methods Europe on Formal Methods for Increasing Software Productivity*, in FME '01. Berlin, Heidelberg: Springer-Verlag, 2001, pp. 500–517.

References

- [8] A. R. Bradley, “SAT-based model checking without unrolling,” in *Proceedings of the 12th International Conference on Verification, Model Checking, and Abstract Interpretation*, in VMCAI'11. Austin, TX, USA: Springer-Verlag, 2011, pp. 70–87.
- [9] S. Dinesh, Y. Zhu, and C. W. Fletcher, “H-Houdini: Scalable Invariant Learning,” in *Proceedings of the 30th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 1*, 2025, pp. 603–618.
- [10] N. Halbwachs and P. Raymond, “A tutorial of lustre,” *Lustre V4*, January, 2002, [Online]. Available: https://www-verimag.imag.fr/DIST-TOOLS/SYNCHRONE/lustre-v4/distrib/lustre_tutorial.pdf

Initial fiddling with SMT solvers: encoding programs

Two encodings of the same Lustre program:

```
(define-fun-rec x_naive ((n Int)) Int
  (ite (= n 0) 1 (+ (x_naive (- n 1)) 1)))
(define-fun-rec y_naive ((n Int)) Int
  (ite (= n 0) 1 (+ (y_naive (- n 1)) (x_naive (- n 1)))))
(define-fun ok_naive ((n Int)) Bool
  (>= (y_naive n) 0))

(declare-const n Int)

#echo "consecution: unsat iff ok(n) => ok(n+1) is true:"
(assert (and (>= n 0) (ok n) (not (ok (+ n 1)))))
(check-sat)
```

Z3 **timeouts** without providing a counterexample.

```
x = 1 -> pre x + 1;
y = 1 -> pre (x + y);
ok = y >= 0;
```

```
(define-fun init ((x Int) (y Int)) Bool (and (= x 1) (= y 1))) smt2

(define-fun trans ((x Int) (y Int) (nx Int) (ny Int)) Bool
  (and (= nx (+ x 1))
        (= ny (+ y x))))
(define-fun ok ((y Int)) Bool (>= y 0))

(declare-const x Int)
(declare-const y Int)
(declare-const nx Int)
(declare-const ny Int)

#echo "consecution: unsat iff ok is inductive:"
(assert (and (ok y)
             (trans x y nx ny)
             (not (ok ny))))
(check-sat)
```

Z3 **immediately answers** SAT, providing a counterexample.

Lustre \Rightarrow SMT: tuples

Two possible solutions:

- define tuple sorts directly in Z3's logic
 - **treat a n -tuple as n expressions**: return a list of expr defining each member of the tuple.
-

```
let rec compile_expr
  (ctx    : context   )
  (env    : z3_env_t  )
  (n      : Expr.expr)
  (n_pre : int       )
  (n_arr : int       ) : Expr.expr list = ...
```

Ocaml

Remarks:

- nested tuples are flattened
- we considered **if** to be the only polymorphic operator [10]
(other operators such as $=$, $+$, ... do not support tuples)

Lustre \Rightarrow SMT: function calls

Problem:

```
node incr(x: int) returns (y: int);  
var l: int;  
let  
    l = x + 1;  
    y = x + l;  
tel  
  
node compute(a, b: int) returns (c: int);  
var z, t: int;  
let  
    z = incr(a);  
    t = incr(b);  
    b = z + t;  
tel
```

lustre

We can't just have

$$(\text{incr def}) \quad \begin{cases} l(n) = x(n) + 1 \\ y(n) = x(n) + l(n) \end{cases}$$

$$(\text{incr calls}) \quad \begin{cases} x(n) = a(n) \\ z(n) = y(n) \\ \textcolor{red}{x(n) = b(n)} \\ \textcolor{red}{t(n) = y(n)} \end{cases}$$

We need to resort to **instantiating** the node at each call site (*i.e.* perform inlining).

(And that was a very painful rabbit hole, because I wanted definitions to be functions of $n: \text{expr} \rightarrow \text{expr}$, instead: use substitutions.)

Lustre \Rightarrow SMT: function calls

$$\text{(incr call 1)} \quad \begin{cases} x_a(n) = a(n) \\ l_a(n) = x_a(n) \\ y_a(n) = x_a(n) + l_a(n) \\ z(n) = y_a(n) \end{cases}$$

$$\text{(incr call 2)} \quad \begin{cases} x_b(n) = b(n) \\ l_b(n) = x_b(n) \\ y_b(n) = x_b(n) + l_b(n) \\ t(n) = y_b(n) \end{cases}$$

Problem: we now have to learn invariant for **each instance** (separately in the worst case).

How to learn efficiently invariants for a node (and not just for its instances)?

In real-world projects (see [discussion #1256](#) and [5], [6]):

- compositional reasoning
- modular reasoning
- progressive refinements
- ...

This issue will **not** be addressed in the project.

Positive examples for transition system

Strategy: $(C \wedge T \Rightarrow C') \Leftrightarrow C \wedge T \wedge \neg C'$ to see if C is inductive (iff query UNSAT).

Problem: **if C is UNSAT alone (e.g. $C = x \geq 0 \wedge x < 0$), we **cannot refine** C any further.**

\Rightarrow [6], [9]: use **positive examples** (concretely: some traces of the program) to first sift C

Problem: initial state alone is not sufficient: **pre** variables are not initialized

- $x = 0 \rightarrow 1 \rightarrow \text{pre pre } x$
- $x = \text{ite}(S_0^\rightarrow, 0, \text{ite}(S_1^\rightarrow, 1, S_1^{\text{pre}}))$
 - $S_1^{\text{pre}'} = S_0^{\text{pre}}$
 - $S_0^{\text{pre}'} = x$

t	x	S_0^\rightarrow	S_1^\rightarrow	S_1^{pre}	S_0^{pre}
0	0	true	false	\emptyset	\emptyset
1	1	false	true	\emptyset	0
2	0	false	false	0	1

\Rightarrow simulate the program for k steps (where k denotes the max depth of **pre** statements):
difficult with this encoding